Executive Summary

AI-FinOps: LLM Cost, Performance & Optimization Analytics

Purva Bachal
GitHub: View GitHub Repository | **Tableau Public:** View Published Dashboards

Tech Stack: Python | Pandas | PostgreSQL | Tableau Public

Bottom Line

Across the **5-7 October 2026 analysis cycle**, the model catalog stayed in the mid-460s while availability and benchmark coverage changed enough to make a static cheapest-model ranking unreliable. The **6 October 2026 dashboards** therefore evaluate six workloads using capability fit and cost together, with the 7 October refresh confirming **465 current models** and a **62.44% benchmark-key match rate**.

The Business Context

The challenge: LLM adoption creates a model-selection problem: teams need enough capability for the workload without paying for more capability than they need.

The approach: Compare model pricing and benchmark-based capability across six representative workload profiles to evaluate cost, capability and trade-offs.

Key Findings

- **The market was active, not static:** The 5 October pricing baseline contained **466 models**; the 7 October refresh contained **465 current models** and recorded **2 new and 1 removed model**. The 6 October dashboard should therefore be read as a dated decision view within a changing catalog.
- **Price concentration is high, but price is not the decision rule:** On 5 October, **437 models (93.78%)** were priced, 22 (**4.72%**) were free and 7 (**1.50%**) were pricing exceptions. The six-workload analysis shows that a low-cost or free option still needs to meet the relevant capability requirement.
- **Capability evidence is incomplete and workload-specific:** The 7 October refresh recorded a **62.44% benchmark-key match rate**. Across the six workloads analysed on 6 October, the relevant capability metric and token assumptions differed by use case, so the best practical option changes with the workload.

- **The decision is a balance of cost, capability and evidence:** The useful shortlist is not the cheapest set of models; it is the set that meets the workload requirement at a sensible cost with enough comparable evidence to support the choice.

Recommendations

- **Set a capability threshold for each workload, then optimize cost** among models that qualify.
- **Use lower-cost qualifying models as the default for routine work** and reserve higher-cost, higher-capability options for workloads where the additional capability is justified.
- **Use benchmark coverage and pricing status as confidence checks** before making a recommendation.
- **Use the cost-capability frontier as a shortlist, not a permanent winner**, and refresh the decision when the catalog or pricing changes.
- **Keep workload assumptions consistent across refreshes** so market movement is not confused with a change in the analytical setup.

Data Caveats

The published Tableau dashboards reflect the **6 October 2026 analysis snapshot**, while the 5 October baseline and 7 October refresh provide the surrounding market context. The pipeline is manually refreshable, benchmark coverage is incomplete, and benchmark scores are source- and time-specific indicators.

Published Analysis

The published Tableau dashboards use the stable CSV reporting layer generated for the **6 October 2026 analysis**.